

- 1) Start from the goal, not the code
Goal: Model Human factors (training, shift length, fatigue, staffing) so they change how fast alarms are handled and how much downtime accrues.
That implies two effects:
 - A site-level effect that scales reaction/repair times.
 - A per-equipment effect that runs acknowledge/repair timers whenever an alarm occurs.
- 2) Split responsibilities (ECS-style)
ECS rule: state lives in components; behavior in systems.
 - Site/state (applies to everyone): "HumanFactors" -> stores training, fatigues, shift, staff, and outputs a computed multiplier "reaction_time_mult".
 - Pre-asset/state (only on alarmable things): "AlarmResponse" -> stores whether a response is active, ack/repair timers, and the target delays.
 - Org metrics/state (executive KPIs): "SiteKPI" -> counts alarms, tracks downtime seconds.Why three components?
Different lifetimes & scope:
 - "HumanFactors": singleton (one per site)
 - AlarmResponse: Per entity that can alarm.
 - SiteKPI: singleton collector for dashboard.
- 3) Decide data flow (dependencies -> tick order)
List what each stage needs and produces:
 - AlarmSystem: needs plant physics already computed; sets "Alarmable.hi/lo/latched".
 - HumanFactorsSystem: reads "HumanFactors" inputs; produces "reaction_time_mult".
 - ResponseSystem: reads "Alarmable" + "reaction_time_mult"; advances "AlarmResponse" (ack/repair timers), may clear alarm; updates "SiteKPI" (downtime, counts).The Order
-> Hydraulics/Thermal -> AlarmSystem -> HumanFactorSystems -> ResponseSystem -> Analytics/Render
Rationale:
 - You can't evaluate alarms until physics ran.
 - You can't scale response times until human factors multiplier is known.
 - You can't update KPIs until alarms/response were processed.
- 4) Choose where to instantiate
You need at least:
 - One equipment entity that already has "Alarmable" -> add "AlarmResponse" to it.
 - One site entity -> add "HumanFactors" and "SiteKPI" to it.Why site entity not globals?
ECS works best when everything is data in the registry - easy to save / serialize, easy to replace later with a JSON loader or multiple files.
- 5) Pick stable defaults (determinism)
Simulations should run even with sparse configs. So:
 - Give safe defaults in components (e.g., training 0.7, shift 8.0h).
 - Pick demo-friendly response times (shorter ack/repair) so you can see it work immediately.
 - Clamp multipliers to sane bounds (e.g., 0.5 ... 2.0) to avoid runaway timers.

- 6) Keep coupling low (testability & reuse)
 - HumanFactorsSystem: only computes a multiplier. It doesn't poke timers.
 - ResponseSystem: only consumes the multiplier + alarms and updates "AlarmResponse" / "SiteKPI". This separation makes each system easy to test and swap (e.g., try a different fatigue formula).
- 7) UI/KPI tap points (observability)

Executives need numbers, so we expose read-only values in the UI:

 - Tank level (you already had), pumps running (quick health check).
 - SiteKPI: downtime minutes, active alarms.
 - HumanFactors: reaction-time multiplier (so sliders later will "feel" immediate). These are probes - they don't change sim state, they just show it.
- 8) Pattern you can reuse for any new feature

Every new feature follows the same recipe:

 - 1) Responsibility: what does it change each tick?
 - 2) Inputs/outputs: which components does it read/write?
 - 3) Components: add state you must persist.
 - 4) System: put behaviour after its inputs and before its consumers.
 - 5) Instantiate: attach components in scenario/loader.
 - 6) Probe: add KPIs/plots so you can tune it.
- 9) Why those specific names/fields?
 - HumanFactors: fields match the executive levers they'll want on sliders: training, fatigue, shift length, staff. The multiplier is a clean interface for others.
 - AlarmResponse mirrors how plants work: acknowledge -> repair -> clear. Keeping both target times and live timers lets you model scaling and partial progress.
 - SiteKPI: isolates rollups (alarms raised/active, downtime) so analytics stays simple and serializable.
- 10) Edge cases accounted for
 - Transient alarms: if condition clears quickly before ack, we allow auto-clear (prevents sticky responses on blips).
 - No site factors present: default multiplier 1.0 (system still works).
 - Multiple alarmable assets: view iterates all entities; KPIs aggregate naturally.

That's the reasoning chain.